

## 1 EVB for information states

2 We are ultimately interested in how the value of our current information state  $z_{\text{agent}} = \langle s_{\text{agent}}, b_{\text{agent}} \rangle$  changes  
3 as a result of a policy update at some future information state:

$$\text{EVB}(z_k, a_k) = \underbrace{\sum_{z \in \mathcal{Z}} \sum_{i=0}^{\infty} \gamma^i P(z_{\text{agent}} \rightarrow z, i, \pi_{\text{old}})}_{\text{Need}} \times \underbrace{\sum_a [\pi_{\text{new}}(z, a) - \pi_{\text{old}}(z, a)] q_{\pi_{\text{new}}}(z, a)}_{\text{Gain}} \quad (1)$$

4 Note that M&D get rid of the first summation (in their case, it was over  $\mathcal{S}$ ), since they assume that a  
5 single update changes the policy only at that update location (and hence Gain is 0 everywhere except that  
6 location). In our case, this is not quite true, because there are a number of information states that can  
7 potentially benefit from a policy update elsewhere – and thus ideally we want to keep that sum.

8 Our current model has a number of assumptions. For the updates of distal information states, we have:

$$\text{EVB}_{\text{dist}}(z_k = \langle s_k, b_k \rangle, a_k) = \gamma^i P(\langle s_{\rho}, b_{\text{agent}} \rangle \rightarrow \langle s_k, b_k \rangle, i, \pi_{\text{old}}) \times \text{Gain}(\langle s_k, b_k \rangle, a_k) \quad (2)$$

9 where  $\langle s_{\rho}, b_{\text{agent}} \rangle$  is the information state at the root of the state-action tree which contains information  
10 state  $z_k = \langle s_k, b_k \rangle$ . The assumption is therefore that the agent can ‘teleport’ from its current physical state  
11  $s_{\text{agent}}$  to the physical root state of that tree  $s_{\rho}$  without changing its belief  $b_{\text{agent}}$ .

12 As for the root updates, we have the following:

$$\text{EVB}_{\text{root}}(z_{\rho} = \langle s_{\rho}, b_{\rho} \rangle, a_k) = \underbrace{\sum_{i=0}^{\infty} \gamma^i P(\langle s_{\text{agent}}, b_{\text{agent}} \rangle \rightarrow \langle s_{\rho}, b_{\text{agent}} \rangle, i, \pi_{\text{old}})}_{\text{certainty-equivalent Need}} \times \underbrace{\text{Gain}_{\text{sweep}}(\langle s_{\rho}, b_{\text{agent}} \rangle, a_k)}_{\text{swept Gain}} \quad (3)$$

13 Where we partially account for the (fixed horizon) learnt information by our sweeping procedure when  
14 estimating Gain (by using equation 2). Partially because this isn’t reflected in the Need term, since we do  
15 not account for the changes in the transition model when calculating Need.

16 Another limitation is the lack of generalisation across information states because we exclude the summation  
17 over  $\mathcal{Z}$  in equation 1. Mattar & Daw do it as well (in their case the sum is over  $\mathcal{S}$ ) since they assume greedy  
18 updates which means that Gain is non-zero only for the updated experience  $(s_k, a_k)$ .

19 The certainty-equivalence assumption right from the agent’s current physical location  $s_{\text{agent}}$  seems a bit odd,  
20 since we learn something about the transition model and then simply ignore it. My suggestion is to change  
21 how we compute Need to account for at least a fixed-horizon learning, and then do the certainty-equivalence  
22 approximation. This would be more consistent with how we calculate Gain, and in fact then we can have a  
23 single equation for both distal and root updates.

24 The idea is essentially to have two forms of certainty-equivalent Need. The first one would be a non-  
25 cumulative Need of reaching some root state  $s_{\rho}$  from the agent’s current physical location  $s_{\text{agent}}$ . Then, from  
26 that root state we plan up to some fixed horizon, and finish off Need with certainty-equivalent Need based  
27 on the belief reached during that planning. Formally, this would look like this:

$$\begin{aligned} \text{Need}(z_k = \langle s_k, b_k \rangle) &= \sum_{i=0}^N \gamma^i P_{\text{old}}(\langle s_{\text{agent}}, b_{\text{agent}} \rangle \rightarrow \langle s_k, b_{\text{agent}} \rangle, i, \pi_{\text{old}}) && \text{get to the root, fixed belief} \\ &+ \sum_{j=N}^{N+H} \gamma^j P_{\text{upd}}(\langle s_{\text{agent}}, b_{\text{agent}} \rangle \rightarrow \langle s_k, b_k \rangle, j, \pi_{\text{old}}) && \text{consecutive belief updates} \\ &+ \sum_{k=N+H+1}^{\infty} \gamma^k P_{\text{new}}(\langle s_{\text{agent}}, b_k \rangle \rightarrow \langle s_k, b_k \rangle, k, \pi_{\text{old}}) && \text{fixed, updated belief} \end{aligned}$$

28 where  $N$  is the expected number of steps needed to reach the root state  $s_{\rho}$  from the agent’s current location  
29  $s_{\text{agent}}$  and  $H$  is the planning horizon.

- 30 – Might not be worth it depending on the discount  $\gamma$  and/or the distance, since the contributions of  
31 future learning fade away
- 32 – Makes me think – maybe we can derive the approximation error due to a limited horizon analytically?  
33 E.g., some notion of 'information regret'